

1)

pgAdmin 4 Object Tools Edit View Window Help

Object Explorer

- Casts
- Catalogs
- Event Triggers
- Extensions
- Foreign Data Wrappers
- Languages (1)
- Publications
- Schemas (1)
 - public
 - Aggregates
 - Collations
 - Domains
 - FTS Configurations
 - FTS Dictionaries
 - FTS Parsers
 - FTS Templates
 - Foreign Tables
 - Functions
 - Materialized Views
 - Operators
 - Procedures
 - Sequences
 - Tables (1)
 - people
 - Trigger Functions
 - Types
 - Views
- Subscriptions
- postgres
- Login/Group Roles (16)
- Tablespaces (2)

Dashboard Properties SQL Statistics Dependencies Dependents Processes CAP Script.sql CAP/postgres@Po...

Query Query History

```
-- SQL statements for displaying this example data:  
select *  
from People;  
  
select *  
from Customers;  
  
select *  
from Agents;
```

Data Output Messages Notifications

Showing rows: 1 to 9 Page No: 1 of 1

pid [PK] integer	prefix text	firstname text	lastname text	suffix text	homecity text	dob date
1	1	Mr.	Billy	Joel	Piano Man	Oyster Bay
2	2	Ms.	Renee	Rosnes	[null]	Regina
3	3	Sir	Elton	John	Esq.	Pinner
4	4	Mr.	Reginald	Dwight	Pinner	1947-03-25
5	5	Mr.	Michael	McDonald	[null]	St. Louis
6	6	Mr.	Ray	Charles	MD	in Georgia
7	7	Dr.	Stevie	Wonder	Ph.D.	Saginaw
8	8	Ms.	Yuja	Wang (王羽佳)		Beijing
9	10	Dr. (Hon)	Diana	Krall		Nanaimo

Total rows: 9 Query complete 00:00:00.180 LF Ln 135, Col 13

pgAdmin 4 Object Tools Edit View Window Help

Object Explorer

- Casts
- Catalogs
- Event Triggers
- Extensions
- Foreign Data Wrappers
- Languages (1)
- Publications
- Schemas (1)
 - public
 - Aggregates
 - Collations
 - Domains
 - FTS Configurations
 - FTS Dictionaries
 - FTS Parsers
 - FTS Templates
 - Foreign Tables
 - Functions
 - Materialized Views
 - Operators
 - Procedures
 - Sequences
 - Tables (1)
 - people
 - Trigger Functions
 - Types
 - Views
- Subscriptions
- postgres
- Login/Group Roles (16)
- Tablespaces (2)

Dashboard Properties SQL Statistics Dependencies Dependents Processes CAP Script.sql CAP/postgres@Po...

Query Query History

```
from People;  
  
select *  
from Customers;  
  
select *  
from Agents;  
  
select *  
from Products;
```

Data Output Messages Notifications

Showing rows: 1 to 5 Page No: 1 of 1

pid [PK] integer	paymentterms text	discountpct numeric (5,2)
1	1 Net 30	21.12
2	4 Net 15	2.47
3	5 In Advance	5.05
4	7 On Receipt	2.00
5	10 Net 30	10.01

Total rows: 5 Query complete 00:00:00.161 LF Ln 138, Col 16

pgAdmin 4 Object Tools Edit View Window Help

pgAdmin 4

Dashboard x Properties x SQL x Statistics x Dependencies x Dependents x Processes x CAP Script.sql x CAP/postgres@Po...

Object Explorer

- CAP
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wrappers
 - Languages (1)
 - Publications
 - Schemas (1)
 - public
 - Aggregates
 - Collations
 - Domains
 - FTS Configurations
 - FTS Dictionaries
 - FTS Parsers
 - FTS Templates
 - Foreign Tables
 - Functions
 - Materialized Views
 - Operators
 - Procedures
 - Sequences
 - Tables (1)
 - Trigger Functions
 - Types
 - Views
 - Subscriptions
 - postgres
 - Login/Group Roles (16)
 - Tablespaces (2)

Query Query History

```
138 from Customers;
139
140 select *
141 from Agents;
142
143 select *
144 from Products;
145
146 select *
147 from Orders;
148
```

Data Output Messages Notifications

pid	paymentterms	commissionpct
[PK] integer	text	numeric (5,2)
1	2 Quarterly	5.00
2	3 Annually	10.00
3	6 Monthly	1.00
4	7 Weekly	2.00

Total rows: 4 Query complete 00:00:00.145

LF Ln 141, Col 13

pgAdmin 4 Object Tools Edit View Window Help

pgAdmin 4

Dashboard x Properties x SQL x Statistics x Dependencies x Dependents x Processes x CAP Script.sql x CAP/postgres@Po...

Object Explorer

- CAP
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wrappers
 - Languages (1)
 - Publications
 - Schemas (1)
 - public
 - Aggregates
 - Collations
 - Domains
 - FTS Configurations
 - FTS Dictionaries
 - FTS Parsers
 - FTS Templates
 - Foreign Tables
 - Functions
 - Materialized Views
 - Operators
 - Procedures
 - Sequences
 - Tables (1)
 - Trigger Functions
 - Types
 - Views
 - Subscriptions
 - postgres
 - Login/Group Roles (16)
 - Tablespaces (2)

Query Query History

```
138 from Customers;
139
140 select *
141 from Agents;
142
143 select *
144 from Products;
145
146 select *
147 from Orders;
148
```

Data Output Messages Notifications

prodid	name	city	qtyonhand	priceusd
[PK] text	text	text	integer	numeric (10,2)
1	p01	Kurzweil PC2R	Dallas	47 67.76
2	p02	Yamaha CP-80	Newark	2399 51.50
3	p03	Apple //+	Duluth	1979 65.02
4	p04	LCARS module	Duluth	3 17.01
5	p05	Roland 808	Dallas	8675309 16.61
6	p06	PDP-11 operator panel	Beijing	88 88.00
7	p07	Flux Capacitor	Newark	1007 1.00
8	p08	HAL 9000 memory ch...	Newark	200 1.25
9	p09	Oberheim OB-Xa	Regina	1 37900.42

Total rows: 9 Query complete 00:00:00.146

LF Ln 144, Col 15

pgAdmin 4

Object Explorer

- CAP
 - Aggregates
 - Collations
 - Domains
 - FTS Configurations
 - FTS Dictionaries
 - FTS Parsers
 - FTS Templates
 - Foreign Tables
 - Functions
 - Materialized Views
 - Operators
 - Procedures
 - Sequences
 - Tables (1)
 - orders
 - Trigger Functions
 - Types
 - Views
 - Subscriptions
- postgres
 - Login/Group Roles (16)
 - Tablespaces (2)

Query Editor

Query: CAP/postgres@PostgreSQL 17

```

140 select *
141 from Agents;
142
143 select *
144 from Products;
145
146 select *
147 from Orders;
  
```

Data Output

Showing rows: 1 to 14 | Page No: 1 of 1

orderid	dateordered	custid	agentid	prodid	quantityordered	totalsales
1	2024-01-22	1	2	p01	1100	58794.00
2	2023-01-23	4	3	p03	1200	76096.81
3	2022-01-23	5	3	p05	1000	15771.20
4	2021-01-23	7	3	p01	1000	66404.80
5	2023-02-14	1	3	p03	500	25643.98
6	2023-02-14	1	3	p04	600	8050.49
7	2023-02-14	1	2	p02	400	16249.28
8	2023-02-14	4	6	p07	600	585.18
9	2023-02-14	4	6	p01	1000	66086.33
10	2023-03-15	1	3	p06	450	31236.48
11	2023-03-15	1	2	p05	500	6550.98
12	2023-03-15	5	2	p01	880	56671.55
13	2022-04-01	7	3	p07	888	870.24
14	2022-05-04	7	6	p03	808	47277.29

Total rows: 14 | Query complete 00:00:00.182

2. Explain the distinctions among the terms primary key, candidate key, and superkey.

These particular keys go down in order, from super key, to candidate key, then to primary key. Super keys are any column / set of columns that uniquely identify every row in a given table. This is the most broad key, in that it will take in any attribute that uniquely identifies a row, so in the CAP people table for example, it allows pid, first name, last name, and dob to be super keys. Candidate keys will narrow down these super keys, as they are defined to be a minimal super key. This allows us to take away first and last name and keep pid and dob, as these attributes have more of a chance to be unique identifiers, as names can certainly repeat. Primary keys narrow down the candidate key to just one, so in this case, we would choose pid since it is guaranteed to be different in each row; making this the most specific key out of these three.

3. Write a short essay on data types. Select a topic for which you might create a table. Name the table and list its fields (columns). For each field, give its data type and whether or not it is nullable.

As mentioned in class, many data types that we already know of exist in SQL, such as integers and "text," which upon some quick research, is a variation of a string. We also mentioned in class that there will be some new data types that we haven't seen yet, such as date / timestamps. These data types allow for certain columns to have proper assignment, where one column can only contain one specific data type. Without data types, the data in our databases cannot be properly stored.

Now, let's create an example. Say we are looking at a table of various NBA players, titled just "NBA Players" where the fields include: pid, first name, last name, current team, years on contract, salary, and date of first game played (YYYY-MM-DD). The fields with text as its data type would be first name, last name, and current team. There would be two integer fields with years on contract and pid, one decimal field for their salary (as some salaries may go down to a specific cent), and one date data type for their first game played. For this specific example, I am going to say that each field is not nullable, as every player in the NBA automatically has these associations just by being a player in the league, which will hold true in this table of players.

4. Explain the following relational "rules" with examples and reasons why they are important.

a. The "first normal form" rule

Definition - There can be no multi-valued attributes or values with internal structure at any intersection of a row and column of a table (no repeating groups or fields) -> All values must be **Atomic**, in that they cannot be subdivided (i.e. don't particularly want first and last name in the same column, instead have one for first name and one for last).

This rule is important because a convoluted attribute can be really confusing, and it is just bad design. From the James Bond example in class, the skills column contained multiple as opposed to just one, and this is a violation. Splitting them up into "skill1" and "skill2" simplifies it and fixes this violation, since the multi-valued attribute has been dealt with. However, this can still be improved, where we can create a new table referring to just the skills, and the violation will most definitely have been removed.

Referring to the beloved CAP database, the People table is a great example of a table that does not violate this rule, as there are no repeating groups or fields, and each value is atomic. Each attribute is uniquely defined, and we have a primary key of pid to ensure that each row is uniquely defined with no repeats. Say there happens to be a row with two people that have the same first and last name, same date of birth, same home city, suffix, and prefix. These two rows will still be unique, as their pid's will be different, signifying that they are different people.

b. The "access rows by content only" rule

Definition - We can only ask for (query) data by what is there, never by where it is. This means that we need to be very specific in what we are looking for, as our tables are ultimately sets, and have no intrinsic order. For example (using my NBA table), we cannot ask "what is the current team in the second row?" Instead, we should ask "What is the current for pid 2?"

This rule is important as we (should) want to be as specific as possible when searching what we want to look for. By asking the question that we should ask, we get right to the point, and we don't have to incorrectly think that our table has an intrinsic order. Just by asking for data specifically regarding pid 2, we can get anything we may need from it.

To build off of the above example, the correct way access our rows is the outline for a basic SQL query, so translate to it, we would say "SELECT currentTeam | FROM NBA Players | WHERE pid = 2;" overall demonstrating that this language properly follows this rule, enforcing you to create your queries the correct way (at least in Postgres).

c. The “all rows must be unique” rule

The definition is practically in the name, in that every row truly must be unique. If two rows were identical, we would not be able to identify their differences.

This is why something like a pid, or in our Marist case, a CWID is necessary, in that it establishes uniqueness, and would allow us to see what subtypes a certain person or other object is part of. For example, in the CAP database, we can see that Stevie Wonder is both a customer and an agent. Without the use of pid, we probably would not be able to point that out, and it can create confusion in that we could think that the customer and agent are two different people. Using the same example from part a, uniqueness is super necessary with the use of an ID. As mentioned, we could have two people with the same first name, last name, DOB, and so on, but because an ID is present, the rows automatically become unique (further demonstrating the importance of primary keys).

Overall, this row cannot be overlooked, as a duplicate row is a big mistake in practice.